

# Pandas

Inteligencia Artificial en los Sistemas de Control Autónomo  
Máster en Ciencia y Tecnología desde el Espacio

Departamento de Automática

## Objectives

1. Introduce Series and DataFrame data structures
2. Understand Pandas features
3. Fluent data manipulation with Pandas
4. Data exploration

## Bibliography

Jake VanderPlas. *Python Data Science Handbook*. Chapter 3. O'Reilly. (Link).

# Table of Contents

1. Introduction
2. The Pandas Series object
3. The Pandas DataFrame object
  - Dataframe concept
  - Constructing DataFrame objects
4. Data indexing and selection
  - Series
  - Loc, iloc and ix
5. Operating on data
  - Overview
  - Missing data
6. Combining datasets
  - `pd.concat()`
  - `pd.merge()`
7. Aggregation in Pandas
  - Grouping in Pandas

# Introduction

A DS/ML workflow needs more features

- Missing data
- Data input
- Operations on groups
- Label columns and rows

Pandas provides all those features, and more


- Pandas = **PAN**el **DA**ta **S**ystem
- Built on NumPy's ndarray
- Provides **dataframes**

Pandas provides two main objects

- Series and DataFrame

# pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



## Convention

```
import numpy as np
import pandas as pd
```

# The Pandas Series object (I)

A **Series** is a one-dimensional array of indexed data

- NumPy arrays indices are implicit (i.e. its position)
- Series indices are explicit, and can be any type

INDEX	VALUES
'a'	0.25
'b'	0.5
'c'	0.75
'd'	0.99

Two attributes

- `values`: ndarray
- `index`: `pd.Index` object

Two indices

- Implicit: Regular index
- Explicit: Custom index

```
data = pd.Series([0.25,
                  0.5, 0.75, 1.0])
data.values
data.index
data[1:3]
```

# The Pandas Series object (II)

```

Custom indices

In [1] : data = pd.Series([0.25, 0.5, 0.75, 1.0],
                          index=['a', 'b', 'c', 'd'])

In [2]: data
Out[1]:
a      0.25
b      0.50
c      0.75
d      1.00
dtype: float64

In [3]: data['a']
Out[2]: 0.25
In [4]: data[0]
Out[3]: 0.25
  
```

# The Pandas DataFrame object

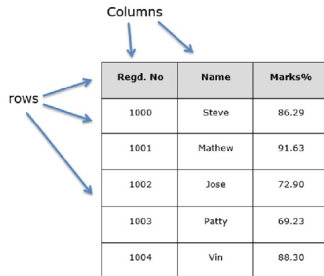
## Dataframe concept (I)

A DataFrame is a 2-D tabular data structure

- Similar to a spreadsheet
- Homogeneous columns
- Heterogeneous rows

Two read-only attributes, both `pd.Index`

- `index`: Rows
- `columns`: Columns



The diagram shows a table with 3 columns and 6 rows. The first row is a header with columns 'Regd. No', 'Name', and 'Marks%'. The subsequent rows contain data for students 1000 through 1004. Annotations include 'Columns' with arrows pointing to the header row, and 'rows' with arrows pointing to the first, second, and last rows of the data section.

Regd. No	Name	Marks%
1000	Steve	86.29
1001	Mathew	91.63
1002	Jose	72.90
1003	Patty	69.23
1004	Vin	88.30

(Source)

# The Pandas DataFrame object

## Dataframe concept (II)

### DataFrame example

```
In [1]: import seaborn as sns
```

```
In [2]: iris = sns.load_dataset('iris')
```

```
In [3]: iris.head()
```

```
Out [1]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	setosa
1	4.9	3.0	1.4	0.2	setosa
2	4.7	3.2	1.3	0.2	setosa
3	4.6	3.1	1.5	0.2	setosa
4	5.0	3.6	1.4	0.2	setosa

```
In [246]: iris.columns
```

```
Out [246]:
```

```
Index(['sepal_length', 'sepal_width', 'petal_length',  
       'petal_width', 'species'], dtype='object')
```



# The Pandas DataFrame object

## Constructing DataFrame objects (I)

### Manual initialization

- From a single Series object  
`pd.DataFrame(population, columns=['population'])`
- From several Series objects  
`pd.DataFrame('population': population, 'area': area)`
- From a dictionary  
`pd.DataFrame([{'a': 0, 'b': 0}, {'a': 1, 'b': 2}])`
- From a NumPy 2-D array  
`pd.DataFrame(np.random.rand(3, 2),  
 columns=['foo', 'bar'], index=['a', 'b', 'c'])`

# The Pandas DataFrame object

## Constructing DataFrame objects (II)

Read from a file

- CSV (very common!!!): `pd.read_csv('filename.csv')`
- Excel:  
`pd.read_excel('filename.xlsx', sheetname='mysheet')`

### CSV example

```
# This CSV file contains data about weights and heights
"id", "weight", "height", "sex", "race"
1, 143.5, 81.6, "Female", "White"
2, 109.1, 83.7, "Female", "Black"
4, 104.8, 54.6, "Female", "Hisp"
7, 130.2, 81.7, "Male", "White"
```

CVS can be exported from MS Excel or programatically

# Data indexing and selection

## Series

### Dictionary-like syntax

```

>> data = pd.Series([0.25, 0.5,
                     0.75, 1.0], index=['a', 'b',
                     'c', 'd'])

>> 'a' in data
True

>> data.keys()
Index(['a', 'b', 'c'], dtype='object')

>> list(data.items())
[( 'a', 0.25), ( 'b', 0.5), ( 'c',
0.75)]

>> data['e'] = 1.25
    
```

### Array-like syntax

```

>> data['a':'c'] # Explicit index
a    0.25
b    0.50
c    0.75
dtype: float64

>> data[0:2] # Implicit index
a    0.25
b    0.50
dtype: float64

>> data[data > 0.5] # Masking
c    0.75
d    1.00
dtype: float64

>> data[['b', 'c']] # Fancy index
b    0.5
c    0.75
dtype: float64
    
```

# Data indexing and selection

## DataFrame

### Dictionary-like syntax

```

>> data['area']
>> data.area
>> data.area is data['area']
True
>> data['density'] = data['pop']
    ]/ data['area']
^^ I ^^ I ^^ I
    
```

### Array-like syntax

```

>> data.values # Get values
    array
>> data.T # Transpose
>> data[0] # First row
>> data['area'] # Area column
    
```

### Remember indexing conventions

- Indexing refers to columns (`data['area']`)
- Slicing refers to rows (`data['Florida':'Illinois']`)
- Masking refers to rows (`data[data.density > 100]`)

# Data indexing and selection

## loc, iloc and ix

### Two types of indices in Pandas

- Explicit and implicit
- Indexing (`data[0]`) is explicit
- Slicing (`data[:2]`) is implicit (Python-like)
- Source of troubles!

Pandas makes explicit the used scheme

- `loc`: Explicit index
- `iloc`: Implicit index
- `ix`: Hybrid

```
# Series
>> serie.loc[1]
>> serie.loc[1:3]
>> serie.iloc[1]
>> serie.iloc[1:3]

# Dataframes
>> df.iloc[:3, :2]
>> df.loc[: 'illinois ', : 'pop ' ]
>> df.ix[:3, : 'pop ' ]
>> df.loc[df.data>100, [ 'pop ', '
    density ' ]]
>> df.iloc[0, 2] = 90
```

# Operating on data

## Overview (I)

Pandas fully supports NumPy's ufuncs

- Efficient computations

Additional Pandas features

- Index and column name preservation
- Index aligning
- Easy data combination

```

>> rng = np.random.RandomState(42)
>> df = pd.DataFrame(rng.randint(0,
                                10, (3,4)))
>> df = pd.DataFrame(rng.randint(0,
                                10, (3,4)), columns=['A', 'B', 'C',
                                'D'])
>> print(df)
   A  B  C  D
0  7  2  5  4
1  1  7  5  1
2  4  0  9  5
>> np.sin(df * np.pi / 4)
   A         B         C         D
0 -7.07e-01  1.0   -0.7   1.22e-16
1  7.07e-01  -0.7  -0.7   7.07e-01
2  1.22e-16  0.0   0.7  -7.07e-01
  
```

# Operating on data

## Overview (II)

### Index preservation

```

>> A = pd.Series([2, 4, 6], index=[0, 1, 2])
>> B = pd.Series([1, 3, 5], index=[1, 2, 3])
>> A + B
0      NaN
1      5.0
2      9.0
3      NaN
dtype: float64
>> A.add(B, fill_value=0)
0      2.0
1      5.0
2      9.0
3      5.0
dtype: float64
  
```

# Operating on data

## Missing data (I)

NumPy supports missing data in floating-point data

- Specific value defined by IEEE
- Available as `np.nan`

Pandas supports missing data through two mechanisms

- `None` object, interpreted as NaN (Not a Number)
- `np.nan`: for floating-point data
- Almost automatic NaN handling (types upcast)

```
>> pd.Series([1, np.nan, 2, None])
0      1.0
1      NaN
2      2.0
3      NaN
dtype: float64
```



# Pandas

## Missing data (II)

### Useful functions for missing data

- `isnull()`: Boolean mask with missing data
- `notnull()`: Opposite of `isnull()`
- `dropna()`: Filtered data
- `fillna()`: NaNs filled

```

>> data = pd.Series([1, np.nan,
                    'hello', None])
>> data[data.notnull()]
0          1
2      hello
dtype: object

>> data.dropna()
0          1
2      hello
dtype: object

>> data.fillna(0)
0          1
1          0
2      hello
3          0
dtype: object
  
```

# Combining datasets

## pd.concat() (I)

Many times we need to combine two or more datasets

- Pandas provides `pd.concat()`, `append()` and `pd.merge()`

### pd.concat() signature

```
pd.concat(objs, axis=0, join='outer', join_axes=None,
          ignore_index=False, keys=None, levels=None, names=
          None, verify_integrity=False, copy=True)
```

By default, `pd.concat()` joins rows preserving index

- `axis`: Join columns (`axis=1`)
- `verify_integrity`: Raise error if duplicates (`verify_integrity=True`)
- `ignore_index`: Create new index (`ignore_index=True`)
- `join`: Can be 'outer' (union) or 'inner' (intersection)

# Combining datasets

## pd.concat() (II)

```
>> df1 = pd.DataFrame([{'A': 'A0', 'B': 'B0'}, {'A': 'A1', 'B': 'B1'}])
>> df2 = pd.DataFrame([{'A': 'A2', 'B': 'B2'}, {'A': 'A3', 'B': 'B3'}])

>> print(df1), print(df2); print(pd.concat([df1, df2]))
  A  B      A  B      A  B
0 A0 B0    0 A2 B2    0 A0 B0
1 A1 B1    1 A3 B3    1 A1 B1
                  0 A2 B2
                  1 A3 B3

>> pd.concat([df1, df2], axis=1)
  A  B  A  B
0 A0 B0 A2 B2
1 A1 B1 A3 B3

>> df1.append(df2)
```

# Combining datasets

`pd.merge()` (I)

Merging based on relational algebra

- Similar to databases tables joins
- Pretty intelligent figuring out the desired output
- By default, join dataframes using shared columns names

# Combining datasets

## pd.merge() (II)

### One-to-one

```
>> print(df1); print(df2)
  employee      group
0      Bob  Accounting
1      Jake  Engineering
2      Lisa  Engineering
3       Sue           HR
  employee  hire_date
0      Lisa      2004
1       Bob      2008
2      Jake      2012
3       Sue      2014
>> print(pd.merge(df1, df2))
  employee  group  hire_date
0      Bob  Accounting      2008
1      Jake  Engineering      2012
2      Lisa  Engineering      2004
3       Sue      HR          2014
```

### Many-to-one

```
>> print(df3); print(df4)
  employee  group  hire_date
0      Bob  Accounting      2008
1      Jake  Engineering      2012
2      Lisa  Engineering      2004
3       Sue           HR      2014
  group  supervisor
0  Accounting      Carly
1  Engineering      Guido
2           HR      Steve
>> print(pd.merge(df3, df4))
  employee  group  hire_date  supervisor
0      Bob  Accounting      2008      Carly
1      Jake  Engineering      2012      Guido
2      Lisa  Engineering      2004      Guido
3       Sue           HR      2014      Steve
```

# Combining datasets

## pd.merge() (III)

### Many-to-many

```
>> print(df1); print(df5)
```

	employee	group		group	skills
0	Bob	Accounting	0	Accounting	math
1	Jake	Engineering	1	Accounting	spreadsheets
2	Lisa	Engineering	2	Engineering	coding
3	Sue	HR	3	Engineering	linux
			4	HR	spreadsheets
			5	HR	organization

```
>> pd.merge(df1, df5)
```

	employee	group	skills
0	Bob	Accounting	math
1	Bob	Accounting	spreadsheets
2	Jake	Engineering	coding
3	Jake	Engineering	linux
4	Lisa	Engineering	coding
5	Lisa	Engineering	linux
6	Sue	HR	spreadsheets
7	Sue	HR	organization

# Combining datasets

## pd.merge() (IV)

### pd.merge() signature

```
pd.merge(left, right, how='inner', on=None, left_on=None,
         right_on=None, left_index=False, right_index=False, sort=
         False, suffixes=('_x', '_y'), copy=True, indicator=False,
         validate=None)
```

### Arguments:

- **on**: Key column name
- **left\_on**: Left table key column name
- **right\_on**: Right table key column name
- **how**: Set arithmetic, 'inner' (default, intersection), 'outer' (union, fills missings with NaNs), 'left' (left entries), 'right' (right entries)

# Combining datasets

pd.merge() (V)

```
>>> A
   lkey  value
0   foo     1
1   bar     2
2   baz     3
3   foo     4

>>> B
   rkey  value
0   foo     5
1   bar     6
2   qux     7
3   bar     8

>>> A.merge(B, left_on='lkey', right_on='rkey', how='outer')
   lkey  value_x  rkey  value_y
0   foo     1    foo     5
1   foo     4    foo     5
2   bar     2    bar     6
3   bar     2    bar     8
4   baz     3   NaN    NaN
5   NaN    NaN   qux     7
```



# Aggregation in Pandas (I)

The first step in data analysis is summarization

- First contact with data
- Insight to the dataset

Aggregation methods

- Applied to columns

AGGREGATION	DESCRIPTION
<code>count()</code>	Total number of items
<code>first(), last()</code>	First and last item
<code>mean(), median()</code>	Mean and median
<code>min(), max()</code>	Minimum and maximum
<code>std(), var()</code>	Standard dev. and variance
<code>mad()</code>	Mean absolute deviation
<code>prod()</code>	Product of all items
<code>sum()</code>	Sum of all items
<code>describe()</code>	Data summary

```
>> import seaborn as sns
>> planets = sns.load_dataset('planets')
>> planets.head()
```

	method	number	orbital_period	mass	distance	year
0	Radial Velocity	1	269.300	7.10	77.40	2006
1	Radial Velocity	1	874.774	2.21	56.95	2008
2	Radial Velocity	1	763.000	2.60	19.84	2011
3	Radial Velocity	1	326.030	19.40	110.62	2007
4	Radial Velocity	1	516.220	10.50	119.47	2009

```
>> planets.dropna().describe()
```

	number	orbital_period	mass	distance	year
count	498.00	498.000000	498.00	498.0000	498.000
mean	1.73	835.778671	2.50	52.0682	2007.377
std	1.17	1469.128259	3.63	46.5960	4.167
min	1.00	1.328300	0.00	1.3500	1989.000
25%	1.00	38.272250	0.21	24.4975	2005.000
50%	1.00	357.000000	1.24	39.9400	2009.000
75%	2.00	999.600000	2.86	59.3325	2011.000
max	6.00	17337.500000	25.00	354.0000	2014.000

```
>> planets.mean()
```

number	1.785507
orbital_period	2002.917596
mass	2.638161
distance	264.069282
year	2009.070531
dtype:	float64

# Grouping in Pandas (I)

Aggregation is generally used ...

- ... good to operate with the whole dataset ...
- ... but also is usually insufficient

We need conditional aggregations

- Aggregate conditionally on some label

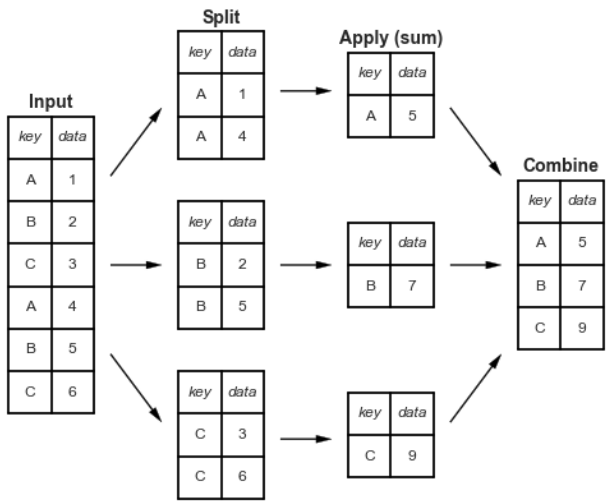
This is done with the operation `groupby` (yes, that name comes from SQL)

- Example: `df.groupby("key")`

Three tasks in one step

1. Split: Break up dependening on a key
2. Apply: Compute some function
3. Combine: Merge results into an output

# Grouping in Pandas (II)





# Grouping in Pandas (IV)

Several mapping methods available

- List  
`df.groupby([2,3,4,1]).sum()`
- Dictionary  
`df.groupby('A': 'vowel', 'B': 'consonant', 'C': 'vowel')`
- Python function  
`df.groupby(str.lower)`
- Multiple keys  
`planets.groupby(['method', 'year'])`
- Mixed keys  
`df.groupby(['key1', 'key2', str.lower])`

# Grouping in Pandas (V)

The method `groupby()` returns an object `groupby`

- Basically, it is a collection of dataframes  
`planets.groupby('method').get_group('Transit')`
- Column selection as dataframe  
`planets.groupby('method')['year']`

Interesting `groupby` attribute, `groups`

- Dictionary with groups  
`planets.groupby('method').groups`
- Compatible with the `len()` method  
`len(planets.groupby('method'))`

# Grouping in Pandas (VI)

## Usual operations with groupings

- Aggregation:  

```
df.groupby('key').aggregate(['min', np.median, max])
df.groupby('key').aggregate('data1': 'min', 'data2': 'max')
```
- Filtering:  

```
planets.groupby('method').filter(lambda x:
x['distance'].mean() > 50.)
```
- Transformation:  

```
df.groupby('key').transform(lambda x: x - x.mean())
```

`Apply()`: Apply arbitrary function and combine results

- Takes a function as argument that takes a DataFrame  

```
planets.groupby("method").apply(lambda x: x / x.sum())
```



## Grouping in Pandas (VII)

### Grouping by decade

```
decade = 10 * (planets['year'] // 10)
decade = decade.astype(str) + 's'
decade.name = 'decade'
planets.groupby(['method', decade])['number'].sum()
        .unstack().fillna(0)
```