

# Supervised learning

Inteligencia Artificial en los Sistemas de Control Autónomo  
Máster en Ciencia y Tecnología desde el Espacio

Departamento de Automática

## Objectives

1. Extend supervised learning algorithms
2. Apply supervised learning to real-world problems

## Bibliography

- Müller, Andreas C., Guido, Sarah. Introduction to Machine Learning with Python. O'Reilly. 2016

All figures have been taken from

[https://github.com/amueller/introduction\\_to\\_ml\\_with\\_python/blob/master/02-supervised-learning.ipynb](https://github.com/amueller/introduction_to_ml_with_python/blob/master/02-supervised-learning.ipynb)

# Table of Contents

## 1. Generalization, overfitting and underfitting

## 2. k-Nearest Neighbors

- k-NN classification
- kNN regression
- Summary

## 3. Linear models

- Ordinary least squares
- Linear regression
- Regularized linear models
- Ridge regression
- Lasso regression
- ElasticNet

- Regularized linear models comparison

- Linear models for classification

- Summary

## 4. Decision Trees

- Summary

## 5. Ensembles of Decision Trees

- Summary

## 6. Support Vector Machines

- Kernelized Support Vector Machines

- Summary

## 7. A

- b


- A: Summary

- ARIMA

# Generalization, overfitting and underfitting

Generalization: accurate predictions on unseen data

- i.e. there is no overfitting neither underfitting
- Depends on model complexity and data variability

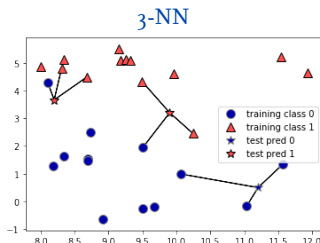
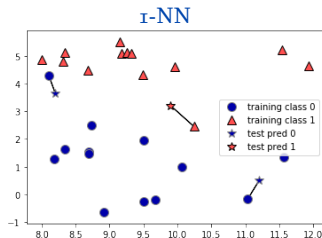


# k-Nearest Neighbors

## k-NN classification (I)

k-NN (k-Nearest Neighbors): Likely, the simplest classifier

- Given a data point, it takes its  $k$  closest neighbors
- Same prediction than its neighbors




k-NN does not generate a model

- The whole dataset must be stored

$k$  uses to be an odd number (1-NN, 3-NN, 5-NN, ...)

# k-Nearest Neighbors

## k-NN classification (II)




$k$  determines the model complexity

- Smoother boundaries in larger  $k$  values
- Model complexity decreases with  $k$
- If  $k$  equals the number of samples,  $k$ -NN always predicts the most frequent class

How to figure out the best  $k$ ?

# k-Nearest Neighbors

## k-NN classification (III)



# k-Nearest Neighbors



## kNN regression (I)

### k-NN regression

Given a data point

1. Take the  $k$  closest data points
2. Predict same target value (1-NN) or average target value (k-NN)


Performance is measured with a regression metric, by default,  $R^2$





# k-Nearest Neighbors

## kNN regression (II)



$k$  determines boundary smoothness

1. With  $k = 1$ , prediction visits all data points
2. With large  $k$  values, fit is worse

# k-Nearest Neighbors

## Summary

Hyperparameters	Advantages	Disadvantages
k	Simple	Slow with large datasets
Distance	Baseline	Bad performance with hundreds or more attributes
		No model
		Dataset must be stored in memory

# Linear models


## Linear model (I)

### Linear model

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_n x_n$$

for a single feature  $y = \beta_0 + \beta_1 x_1$ , where

- $\beta_0$  is the intercept
- $\beta_1$  is the slope
- Interpretable model



Linear models assume a linear relationship among variables

- This limitation can be easily overcome
- Surprisingly good results in high dimensional spaces

# Linear models

## Linear regression

Different linear models for regression

- The difference lies in how  $\beta_i$  parameters are learned

Ordinary Least Squares (OLS): Minimizes mean squared error

- OLS does not have any hyperparameter
- No complexity control

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (x_i - y_i)^2$$

Linear regression can be used to fit non-linear models

- Just adding new attributes

# Linear models

## Regularized linear models

**Regularization:** Term that penalizes complexity

- Added to the cost function
- Linear models remain the same
- Train to minimize cost function and coefficients
- Intercepts are not part of regularization

Three regularizations

- $L_1$  (Lasso regression),  $L_2$  (Ridge regression) and ElasticNet ( $L_1$  and  $L_2$ )

Lasso ( $L_1$ )

$$\alpha \sum_j^n |\beta_j|$$

Ridge ( $L_2$ )

$$\frac{\alpha}{2} \sum_j^n \beta_j^2$$

ElasticNet

$$\alpha \left( \frac{\lambda}{2} \sum_j^n \beta_j^2 + (1 - \lambda) \sum_j^n |\beta_j| \right)$$

# Linear models

## Ridge regression

Ridge regression (or L2 regularization) adds a new term to cost function

$$\text{MSE} + \alpha \sum_{i=1}^n \beta_i^2$$

$\alpha$  controls the model complexity

- If  $\alpha = 0$  Ridge becomes a regular linear regression
- Optimal  $\alpha$  depends on the problem

Ridge by default

# Linear models

## Lasso regression (I)

Lasso regression (or  $L_1$  regularization) adds a new term to cost function

$$\text{MSE} + \alpha \frac{1}{2} \sum_{i=1}^n |\beta_i|$$

$\alpha$  controls the model complexity

- If  $\alpha = 0$  Ridge becomes a regular linear regression
- Optimal  $\alpha$  depends on the problem


Some coefficients may be exactly zero

- Implicit feature selection
- Easier interpretation
- Better with large number of attributes

# Linear models

## Lasso regression (II)

### Dimension Reduction of Feature Space with LASSO



(Source)



# Linear models

## ElasticNet

Lasso and Ridge can be combined

$$\text{MSE} + \alpha \left( \lambda \frac{1}{2} \sum_{i=1}^n |\beta_i| + (1 - \lambda) \sum_{i=1}^n \beta_i^2 \right)$$


Two hyperparameters

- $\alpha$  controls the model complexity
- $\lambda$  balances between  $L_1$  and  $L_2$


# Linear models

## Regularized linear models comparison

### Ridge - L2



### Lasso - L1



# Linear models

## Linear models for classification (I)


A linear regression can be used as classifier

- Just compare the prediction with a threshold (0, for instance)
  - If  $\hat{y} > 0$ , assign class 1
  - If  $\hat{y} \leq 0$ , assign class -1
- The decision boundary for any binary linear classifier is a line, plane or hyperplane

A **logistic regression** is a generalization of a linear regression

- It is a binary classifier
- Its output is a probability



$$\hat{p} = \sigma \left( \beta_0 + \sum_{i=1}^n \beta_i x_i \right), \quad \sigma(x) = \frac{1}{1+e^{-x}}$$



where  $\sigma(t)$  is the logistic function, defined as  $\sigma(t) = \frac{1}{1+e^t}$

# Linear models

## Linear models for classification (II)




# Linear models

## Linear models for classification (III)

The model can be regularized with  $L_1$ ,  $L_2$  and ElasticNet

- In Scikit-Learn, regularization strength is given by  $C$
- Lower values of  $C$  correspond to smaller regularization strength



# Linear models

## Summary

Hyperparameters	Advantages	Disadvantages
-	Fast train and predict	No complexity tuning
$\alpha$ (L1, L2, ElasticNet)	Scales well to large data-sets	Limited in low dimensional spaces
l1_ratio (ElasticNet)	Better in high dimensional spaces	
	Few hyperparameters	
	Interpretable	

Better when the number of features is large compared to the number of samples

# Decission Trees

TODO

# Decission Trees

## Summary

Hyperparameters	Advantages	Disadvantages



# Ensembles of Decision Trees

TODO

# Ensembles of Decision Trees

## Summary

Hyperparameters	Advantages	Disadvantages

# Support Vector Machines

TODO

# Support Vector Machines

## Kernelized Support Vector Machines

TODO

# Support Vector Machines

## Summary

Hyperparameters	Advantages	Disadvantages

A

B

TODO

A

B: Summary

Hyperparameters	Advantages	Disadvantages

# Algorithms

## ARIMA (I)

### AR: Autoregressive model

- Current observation depends on the last  $p$  observations
- Long term memory

AR( $p$ )

$$X_t = c + \sum_{i=1}^p \phi_i X_{t-i} + \epsilon_t$$

### MA: Moving Average model

- Current observation linearly depends on the last  $q$  innovations
- Short term memory

MA( $q$ )

$$X_t = \mu + \epsilon_t + \theta_1 \epsilon_{t-1} + \dots + \theta_q \epsilon_{t-q}$$

ARMA model = AR + MA

- ARMA( $p, q$ ): Two hyperparameters,  $p$  and  $q$




# Algorithms

## ARIMA (II)

ARIMA = AR + i + MA (AR integrated MA)

- ARIMA(p, d, q)
- Three integer parameters: p, q and d (in practice, low order models)

Forecasts from STL + ARIMA(1,1,1) with drift



(Source)

**autoarima:** search over p, q and d