

# Supervised learning

Aprendizaje Automático para la Robótica  
Máster Universitario en Ingeniería Industrial

Departamento de Automática

## Objectives

1. Extend supervised learning algorithms
2. Apply supervised learning to real-world problems

## Bibliography

- Müller, Andreas C., Guido, Sarah. Introduction to Machine Learning with Python. O'Reilly. 2016

All figures have been taken from

[https://github.com/amueller/introduction\\_to\\_ml\\_with\\_python/blob/master/02-supervised-learning.ipynb](https://github.com/amueller/introduction_to_ml_with_python/blob/master/02-supervised-learning.ipynb)

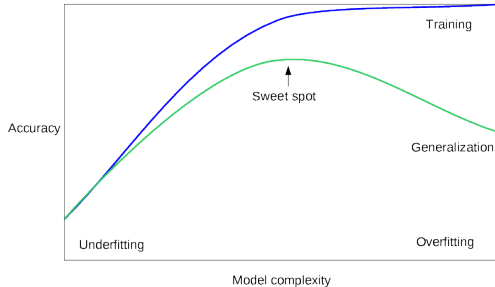
# Table of Contents

1. Generalization, overfitting and underfitting
2. k-Nearest Neighbors
  - k-NN classification
  - Scikit-Learn
  - kNN regression
  - Scikit-Learn
  - Summary
3. Linear models
  - Ordinary least squares
  - Ridge regression
  - Lasso regression
  - ElasticNet
  - Linear models for classification
  - Scikit-Learn
  - Summary
4. Naive Bayes Classifiers
  - Scikit-Learn
  - Summary
5. Decision Trees
  - Scikit-Learn
  - Summary
6. Ensembles of Decision Trees
  - Scikit-Learn
  - Summary
7. Support Vector Machines
  - Kernelized Support Vector Machines
  - Support Vector Machines
  - Summary
8. A
  - b
  - A: Scikit-Learn
  - A: Summary
  - ARIMA

# Generalization, overfitting and underfitting

Generalization: accurate predictions on unseen data

- i.e. there is no overfitting neither underfitting
- Depends on model complexity and data variability



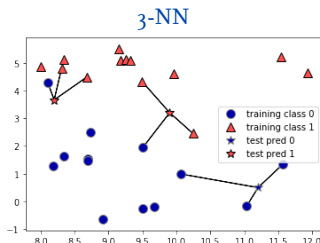
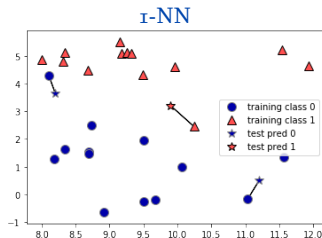
(Source)

# k-Nearest Neighbors

## k-NN classification (I)

k-NN (k-Nearest Neighbors): Likely, the simplest classifier

- Given a data point, it takes its  $k$  closest neighbors
- Same prediction than its neighbors



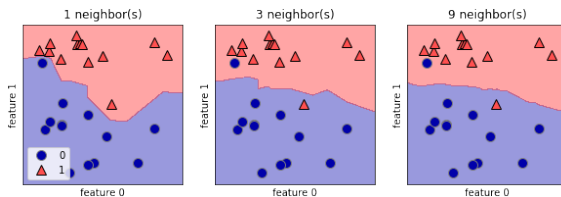
k-NN does not generate a model

- The whole dataset must be stored

$k$  uses to be an odd number (1-NN, 3-NN, 5-NN, ...)

# k-Nearest Neighbors

## k-NN classification (II)



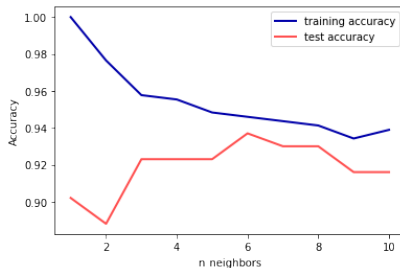
$k$  determines the model complexity

- Smoother boundaries in larger  $k$  values
- Model complexity decreases with  $k$
- If  $k$  equals the number of samples,  $k$ -NN always predicts the most frequent class

How to figure out the best  $k$ ?

# k-Nearest Neighbors

## k-NN classification (III)



# k-Nearest Neighbors classifier

## Scikit-learn

```
sklearn.neighbors.KNeighborsClassifier
```

### Constructor arguments:

- `n_neighbors`: int, default=5
- `metric`: string, default='minkowski'
- `p`: int, default=2 ( $p = 1$  Manhattan distance,  $p = 2$  euclidean distance)

### Attributes:

- `classes_`: ndarray (`n_samples`)

Methods: `fit()`, `predict()`

(Scikit-Learn reference)



# k-Nearest Neighbors

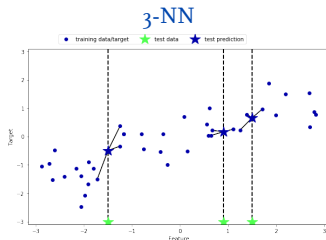
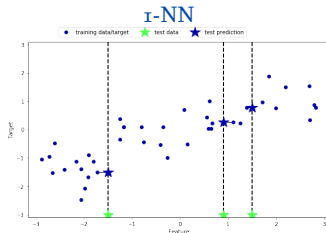
## kNN regression (I)

### k-NN regression

Given a data point

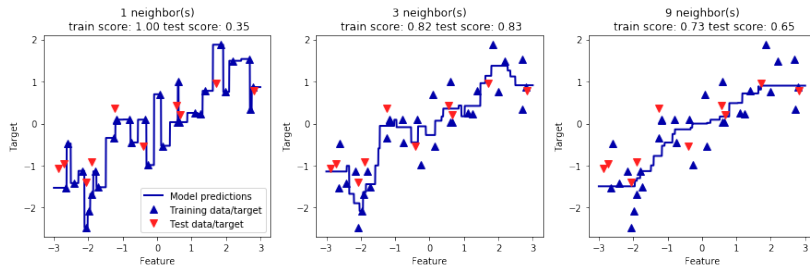
1. Take the  $k$  closest data points
2. Predict same target value (1-NN) or average target value (k-NN)

Performance is measured with a regression metric, by default,  $R^2$



# k-Nearest Neighbors

## kNN regression (II)



$k$  determines boundary smoothness

1. With  $k = 1$ , prediction visits all data points
2. With large  $k$  values, fit is worse

# k-Nearest Neighbors regressor

Scikit-learn

```
sklearn.neighbors.KNeighborsRegressor
```

Constructor arguments:

- `n_neighbors`: int, default=5
- `metric`: string, default='minkowski'
- `p`: int, default=2 ( $p = 1$  Manhattan distance,  $p = 2$  euclidean distance)

Attributes:

Methods: `fit()`, `predict()`

(Scikit-Learn reference)

# k-Nearest Neighbors

## Summary

Hyperparameters	Advantages	Disadvantages
k	Simple	Slow with large datasets
Distance	Baseline	Bad performance with hundreds or more attributes
		No model
		Dataset must be stored in memory

# Linear models

## Linear model (I)

### Linear model

$$y = a_0 + a_1x_1 + a_2x_2 + \dots + a_nx_n$$

for a single feature  $y = a_0 + a_1x_1$ , where

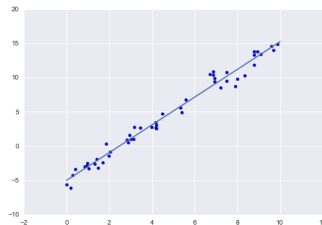
- $a_0$  is the intercept
- $a_1$  is the slope

Linear models assume a linear relationship among variables

- This limitation can be easily overcome
- Surprisingly good results in high dimensional spaces

Different linear models for regression

- The difference lies in how  $a_i$  parameters are learned



## Linear models (II)

Several methods to fit coefficients

- Ordinary Least Squares (OLS)
- Generalized Least Squares (GSL)
- Weighted Least Squares (WLS)
- Generalized Least Squares with AR Covariance Structure (GLSAR)

**Regularization:** Term that penalizes complexity

- $L_1$  (Lasso regression)
- $L_2$  (Ridge regression)
- ElasticNet:  $L_1$  and  $L_2$

Lasso

$$\lambda \sum_j \beta_j^2$$

Ridge

$$\lambda \sum_j |\beta_j|$$

ElasticNet

$$\alpha \sum_j \beta_j^2 + (1 - \alpha) \sum_j |\beta_j|$$

# Linear models

## Scikit-learn

TODO

`sklearn.cluster.AgglomerativeClustering`

Constructor arguments:

- `linkage`: 'ward', 'complete', 'average', 'single'

Attributes:

- `n_clusters`: int
- `labels_`: ndarray (n\_samples)

Methods: `fit()`, `fit_predict()`

(Scikit-Learn reference)

# Linear models

## Summary

Hyperparameters	Advantages	Disadvantages



# Naive Bayes Classifiers

TODO

# Naive Bayes Classifiers

## Scikit-learn

```
sklearn.cluster.AgglomerativeClustering
```

Constructor arguments:

- `linkage`: 'ward', 'complete', 'average', 'single'

Attributes:

- `n_clusters`: int
- `labels_`: ndarray (n\_samples)

Methods: `fit()`, `fit_predict()`

(Scikit-Learn reference)

# Naive Bayes Classifiers

## Summary

Hyperparameters	Advantages	Disadvantages

# Decision Trees

TODO

# Decision Trees

## Scikit-learn

```
sklearn.cluster.AgglomerativeClustering
```

Constructor arguments:

- `linkage`: 'ward', 'complete', 'average', 'single'

Attributes:

- `n_clusters`: int
- `labels_`: ndarray (n\_samples)

Methods: `fit()`, `fit_predict()`

(Scikit-Learn reference)

# Decission Trees

## Summary

---

Hyperparameters	Advantages	Disadvantages
-----------------	------------	---------------

---

# Ensembles of Decision Trees

TODO

# Ensembles of Decision Trees

## Ensembles of Decision Trees : Scikit-learn

```
sklearn.cluster.AgglomerativeClustering
```

Constructor arguments:

- `linkage`: 'ward', 'complete', 'average', 'single'

Attributes:

- `n_clusters`: int
- `labels_`: ndarray (n\_samples)

Methods: `fit()`, `fit_predict()`

(Scikit-Learn reference)



# Ensembles of Decision Trees

## Summary

Hyperparameters	Advantages	Disadvantages

# Support Vector Machines

TODO

# Support Vector Machines

## Kernelized Support Vector Machines

TODO

# Scikit-Learn

# Support Vector Machines

## Scikit-learn

### `sklearn.cluster.AgglomerativeClustering`

#### Constructor arguments:

- `linkage`: 'ward', 'complete', 'average', 'single'

#### Attributes:

- `n_clusters`: int
- `labels_`: ndarray (n\_samples)

Methods: `fit()`, `fit_predict()`

(Scikit-Learn reference)

# Support Vector Machines

## Summary

Hyperparameters	Advantages	Disadvantages

A

B

TODO

A

## B: Scikit-learn

`sklearn.cluster.AgglomerativeClustering`

Constructor arguments:

- `linkage`: 'ward', 'complete', 'average', 'single'

Attributes:

- `n_clusters`: int
- `labels_`: ndarray (n\_samples)

Methods: `fit()`, `fit_predict()`

(Scikit-Learn reference)



A

B: Summary

Hyperparameters	Advantages	Disadvantages

# Algorithms

## ARIMA (I)

### AR: Autoregressive model

- Current observation depends on the last  $p$  observations
- Long term memory

#### AR( $p$ )

$$X_t = c + \sum_{i=1}^p \phi_i X_{t-i} + \epsilon_t$$

### MA: Moving Average model

- Current observation linearly depends on the last  $q$  innovations
- Short term memory

#### MA( $q$ )

$$X_t = \mu + \epsilon_t + \theta_1 \epsilon_{t-1} + \dots + \theta_q \epsilon_{t-q}$$

### ARMA model = AR + MA

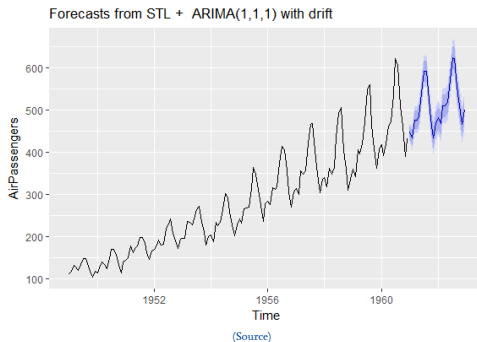
- ARMA( $p, q$ ): Two hyperparameters,  $p$  and  $q$

# Algorithms

## ARIMA (II)

ARIMA = AR + i + MA (AR integrated MA)

- ARIMA(p, d, q)
- Three integer parameters: p, q and d (in practice, low order models)



**autoarima:** search over p, q and d