

Supervised learning

Aprendizaje Automático para la Robótica
Máster Universitario en Ingeniería Industrial

Departamento de Automática

Objectives

1. Extend supervised learning algorithms
2. Apply supervised learning to real-world problems

Bibliography

- Müller, Andreas C., Guido, Sarah. Introduction to Machine Learning with Python. O'Reilly. 2016

All figures have been taken from

https://github.com/amueller/introduction_to_ml_with_python/blob/master/02-supervised-learning.ipynb

Table of Contents

1. Generalization, overfitting and underfitting

2. k-Nearest Neighbors

- k-NN classification
- Scikit-Learn
- kNN regression
- Scikit-Learn
- Summary

3. Linear models

- Ordinary least squares
- Linear regression
- Regularized linear models
- Ridge regression
- Lasso regression
- ElasticNet
- Regularized linear models comparison
- Scikit-Learn
- Linear models for classification
- Scikit-Learn
- Summary

4. Decision Trees

- Building decision trees
- Controlling complexity of decision trees
- Analyzing decision trees
- Feature importance in trees
- Decision trees in regression
- Scikit-Learn
- Summary

5. Ensembles of Decision Trees

- Ensembles
- Random forests
- Analyzing random forests
- Gradient boosted regression trees
- Scikit-Learn
- Summary

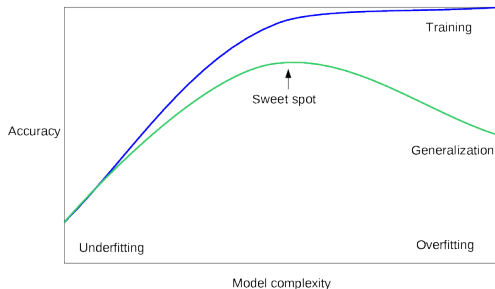
6. Support Vector Machines

- Kernelized Support Vector Machines
- Support Vector Machines
- Summary
- ARIMA

Generalization, overfitting and underfitting

Generalization: accurate predictions on unseen data

- i.e. there is no overfitting neither underfitting
- Depends on model complexity and data variability



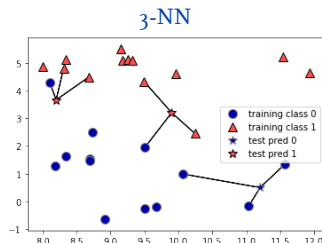
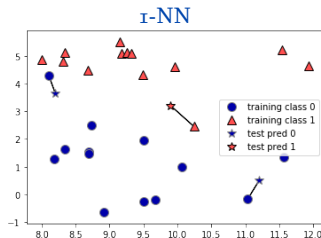
(Source)

k-Nearest Neighbors

k-NN classification (I)

k-NN (k-Nearest Neighbors): Likely, the simplest classifier

- Given a data point, it takes its k closest neighbors
- Same prediction than its neighbors




k-NN does not generate a model

- The whole dataset must be stored

k uses to be an odd number (1-NN, 3-NN, 5-NN, ...)

k-Nearest Neighbors

k-NN classification (II)




k determines the model complexity

- Smoother boundaries in larger k values
- Model complexity decreases with k
- If k equals the number of samples, k -NN always predicts the most frequent class

How to figure out the best k ?

k-Nearest Neighbors

k-NN classification (III)



k-Nearest Neighbors classifier

Scikit-learn

```
sklearn.neighbors.KNeighborsClassifier
```

Constructor arguments:

- `n_neighbors`: int, default=5
- `metric`: string, default='minkowski'
- `p`: int, default=2 ($p = 1$ Manhattan distance, $p = 2$ euclidean distance)

Attributes:

- `classes_`: ndarray(n_{samples})

Methods: `fit()`, `predict()`

(Scikit-Learn reference)

k-Nearest Neighbors

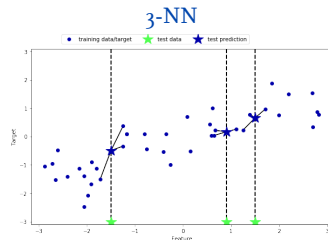
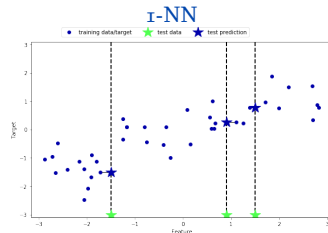
kNN regression (I)

k-NN regression

Given a data point


1. Take the k closest data points
2. Predict same target value (1-NN) or average target value (k-NN)

Performance is measured with a regression metric, by default, R^2



k-Nearest Neighbors

kNN regression (II)



k determines boundary smoothness

1. With $k = 1$, prediction visits all data points
2. With large k values, fit is worse

k-Nearest Neighbors regressor

Scikit-learn

`sklearn.neighbors.KNeighborsRegressor`

Constructor arguments:

- `n_neighbors`: int, default=5
- `metric`: string, default='minkowski'
- `p`: int, default=2 ($p = 1$ Manhattan distance, $p = 2$ euclidean distance)

Attributes:

Methods: `fit()`, `predict()`

(Scikit-Learn reference)

k-Nearest Neighbors

Summary

Hyperparameters	Advantages	Disadvantages
k	Simple	Slow with large datasets
Distance	Baseline	Bad performance with hundreds or more attributes
		No model
		Dataset must be stored in memory

Linear models


Linear model (I)

Linear model

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_n x_n$$

for a single feature $y = \beta_0 + \beta_1 x_1$, where

- β_0 is the intercept
- β_1 is the slope
- Interpretable model



Lineal models assume a linear relationship among variables

- This limitation can be easily overcome
- Surprisingly good results in high dimensional spaces

Linear models

Linear regression

Different linear models for regression

- The difference lies in how β_i parameters are learned

Ordinary Least Squares (OLS): Minimizes mean squared error

- OLS does not have any hyperparameter
- No complexity control

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (x_i - y_i)^2$$

Linear regression can be used to fit non-linear models

- Just adding new attributes

Linear models

Regularized linear models

Regularization: Term that penalizes complexity

- Added to the cost function
- Linear models remain the same
- Train to minimize cost function and coefficients
- Intercepts are not part of regularization

Three regularizations

- L_1 (Lasso regression), L_2 (Ridge regression) and ElasticNet (L_1 and L_2)

Lasso (L_1)

$$\alpha \sum_j^n |\beta_j|$$

Ridge (L_2)

$$\frac{\alpha}{2} \sum_j^n \beta_j^2$$

ElasticNet

$$\alpha \left(\frac{\lambda}{2} \sum_j^n \beta_j^2 + (1 - \lambda) \sum_j^n |\beta_j| \right)$$

Linear models

Ridge regression

Ridge regression (or L2 regularization) adds a new term to cost function

$$\text{MSE} + \alpha \sum_{i=1}^n \beta_i^2$$

α controls the model complexity

- If $\alpha = 0$ Ridge becomes a regular linear regression
- Optimal α depends on the problem

Ridge by default

Linear models

Lasso regression (I)

Lasso regression (or L_1 regularization) adds a new term to cost function

$$\text{MSE} + \alpha \frac{1}{2} \sum_{i=1}^n |\beta_i|$$

α controls the model complexity

- If $\alpha = 0$ Ridge becomes a regular linear regression
- Optimal α depends on the problem

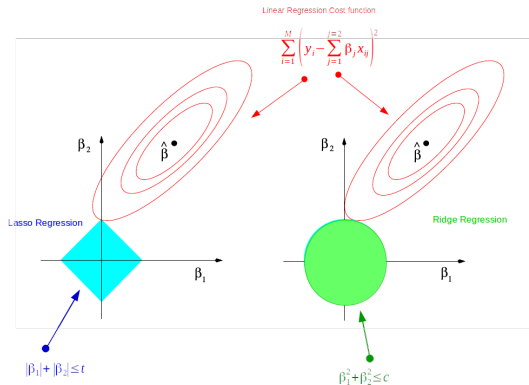
Some coefficients may be exactly zero

- Implicit feature selection
- Easier interpretation
- Better with large number of attributes

Linear models

Lasso regression (II)

Dimension Reduction of Feature Space with LASSO



(Source)

Linear models

ElasticNet

Lasso and Ridge can be combined

$$\text{MSE} + \alpha \left(\lambda \frac{1}{2} \sum_{i=1}^n |\beta_i| + (1 - \lambda) \sum_{i=1}^n \beta_i^2 \right)$$

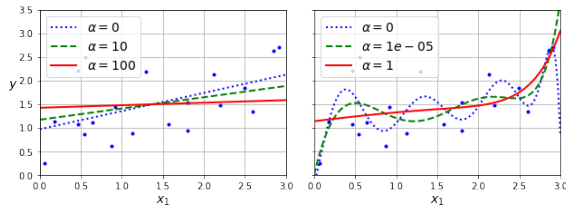
Two hyperparameters

- α controls the model complexity
- λ balances between L1 and L2

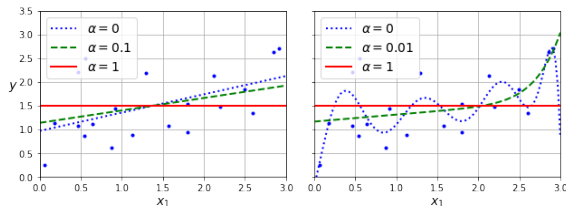
Linear models

Regularized linear models comparison

Ridge - L2



Lasso - L1



Linear models

Scikit-learn (I)

```
sklearn.linear_model.LinearRegression
```

Constructor arguments:

- `fit_intercept`: boolean, default=True

Attributes:

- `coef_`: ndarray (n_features,)
- `intercept_`: ndarray (n_targets,)
- `n_features_in_`: int

Methods: `fit()`, `predict()`

Linear models

Scikit-learn (II)

`sklearn.linear_model.Ridge`

Constructor arguments:

- `fit_intercept`: boolean, default=True
- `alpha`: float, default=1.0

Attributes:

- `coef_`: ndarray (n_features,)
- `intercept_`: ndarray (n_targets,)
- `n_features_in_`: int

Methods: `fit()`, `predict()`

Linear models

Scikit-learn (IV)

```
sklearn.linear_model.ElasticNet
```

Constructor arguments:

- `fit_intercept`: boolean, default=True
- `alpha`: float, default=1.0
- `l1_ratio`: float, default=0.5

Attributes:

- `coef_`: ndarray (n_features,)
- `intercept_`: ndarray (n_targets,)
- `n_features_in_`: int

Methods: `fit()`, `predict()`

Linear models

Scikit-learn (III)

`sklearn.linear_model.Lasso`

Constructor arguments:

- `fit_intercept`: boolean, default=True
- `alpha`: float, default=1.0

Attributes:

- `coef_`: ndarray (n_features,)
- `intercept_`: ndarray (n_targets,)
- `n_features_in_`: int

Methods: `fit()`, `predict()`

Linear models

Linear models for classification (I)

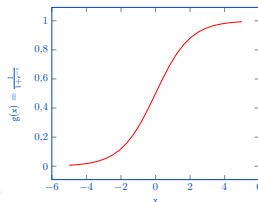
A linear regression can be used as classifier

- Just compare the prediction with a threshold (0, for instance)
 - If $\hat{y} > 0$, assign class 1
 - If $\hat{y} \leq 0$, assign class -1
- The decision boundary for any binary linear classifier is a line, plane or hyperplane

A **logistic regression** is a generalization of a linear regression

- It is a binary classifier
- Its output is a probability

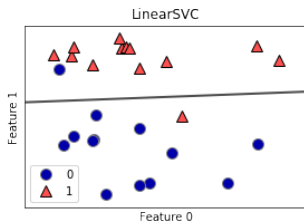
$$\hat{p} = \sigma \left(\beta_0 + \sum_{i=1}^n \beta_i x_i \right), \quad \sigma(t) = \frac{1}{1+e^{-t}}$$



where $\sigma(t)$ is the logistic function, defined as $\sigma(t) = \frac{1}{1+e^{-t}}$

Linear models

Linear models for classification (II)

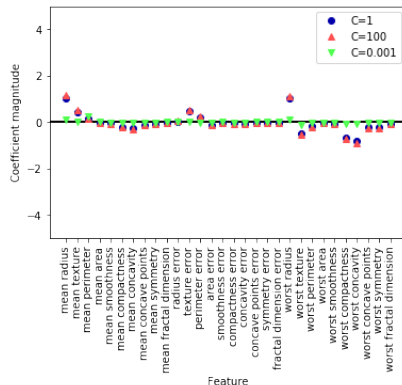


Linear models

Linear models for classification (III)

The model can be regularized with L_1 , L_2 and ElasticNet

- In Scikit-Learn, regularization strength is given by C
- Lower values of C correspond to smaller regularization strength



Linear models

Scikit-learn

`sklearn.linear_model.ElasticNet`

Constructor arguments:

- `penalty`: 'l1', 'l2', 'elasticnet', 'none', default='l2'
- `fit_intercept`: boolean, default=True
- `alpha`: float, default=1.0
- `l1_ratio`: float, default=0.5

Methods: `fit()`, `predict()`

Attributes:

- `coef_`: ndarray (n_features,)
- `intercept_`: ndarray (n_targets,)
- `n_features_in_`: int

Linear models

Summary

Hyperparameters	Advantages	Disadvantages
-	Fast train and predict	No complexity tuning
α (L1, L2, ElasticNet)	Scales well to large data-sets	Limited in low dimensional spaces
l1_ratio (ElasticNet)	Better in high dimensional spaces Few hyperparameters Interpretable	

Better when the number of features is large compared to the number of samples


Decision Trees

Decision trees are a family of algorithms for classification and regression

- They learn a tree data structure
- Hierarchy of if/else questions (test, or node)
- Decision (terminal node or leaf)

Usually, datasets does not contain binary attributes

- Continuous features
- Is feature i larger than value a ?



Decision Trees

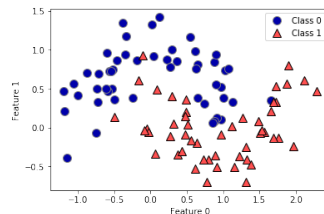
Building decision trees (I)

Tree learning algorithm

1. Begin with the root node
2. Searches all possible tests (according to a purity measure)
3. The most informative test is taken
4. Repeat recursively

Prediction of a new data point

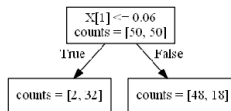
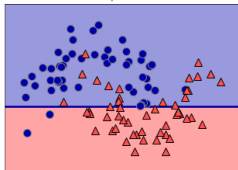
- Classification: Majority class in the partition
- Regression: Average value of target values in the partition



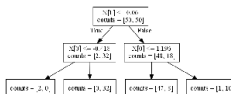
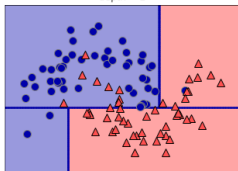
Decision Trees

Building decision trees (II)

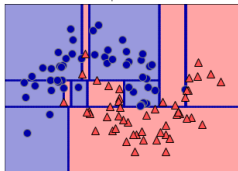
depth = 1



depth = 2



depth = 9



Decision Trees

Building decision trees (III)

Let p_{mk} be the proportion of class k in node m , and Q_m the data in node m

Gini

$$G(Q_m) = \sum_k p_{mk}(1 - p_{mk})$$

Log Loss or Entropy

$$H(Q_m) = - \sum_k p_{mk} \log(p_{mk})$$

Decision Trees

Controlling complexity of decision trees

Trees tend to grow until all leaves are pure

- Very big trees in real problems
- Big trees use to be overfitted models

Two strategies to prevent overfitting

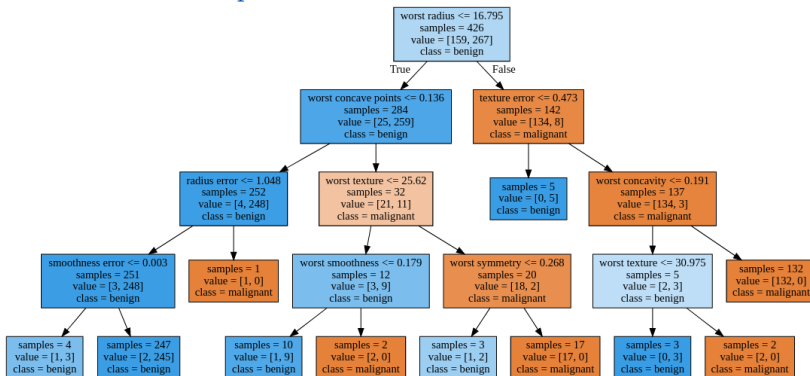
- Pre-pruning: Stop the creation of the tree early according to some criteria
 - Maximum depth, number of leaves, minimum number of points in a node, ...
 - Implemented in Scikit-Learn
- Post-pruning: Build the tree and then remove nodes with little information

Decision Trees

Analyzing decision trees

Decision trees is easily explained to nonexperts

- Interpretable models
- Deep trees are overwhelming
- Trick: Observe the path with most data

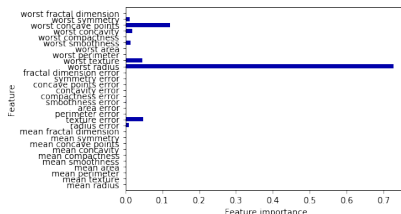


Decision Trees

Analyzing decision trees

Feature importance is a metric that summarizes features

- Number between 0 (not used at all) and 1 (perfect prediction)
- Feature importances sum to one
- Useful for feature selection and model interpretation



Some considerations

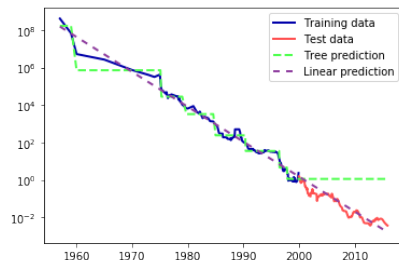
- It does not inform about the relationship between attribute and target
- It quantifies the importance in the tree
 - Correlated attributes may score low importance

Decision Trees

Decision trees in regression

Decision trees are not able to extrapolate

- i. e. to predict outside of the range of the training data
- It is specially important in regression problems



Decision Trees

Scikit-learn (I)

`sklearn.tree.DecisionTreeClassifier`

Constructor arguments:

- `criterion`: 'gini', 'entropy', 'log_loss', default='gini'
- `max_depth`: int, default=None
- `max_leaf_nodes`: int, default=None
- `min_samples_leaf`: int or float, default=1

Attributes:

- `classes_`: ndarray (n_classes,)
- `feature_importances_`: ndarray (n_features,)
- `tree_`: Tree instance

Methods: `fit()`, `predict()`, `decision_path()`, `get_depth()`, `get_n_leaves()`

(Scikit-Learn reference)

Decision Trees

Scikit-learn (II)

`sklearn.tree.DecisionTreeRegressor`

Constructor arguments:

- `criterion`: “squared_error”, “absolute_error”, default=“squared_error”
- `max_depth`: int, default=None
- `max_leaf_nodes`: int, default=None
- `min_samples_leaf`: int or float, default=1

Attributes:

- `feature_importances_`: ndarray (n_features,)
- `tree_`: Tree instance

Methods: `fit()`, `predict()`, `decision_path()`, `get_depth()`, `get_n_leaves()`

(Scikit-Learn reference)

Supervised learning

Decision Trees

Summary

Hyperparameters	Advantages	Disadvantages
max_depth	Visualization	Tend to overfit
max_leaf_nodes	Interpretable by non-experts	Poor generalization
min_samples_leaf	Invariant to scale	
'criterion'	Mix of categorical and numerical data	

Ensembles of Decision Trees


Ensembles

Ensembles, in ML, refers to the combination of several models

- For instance, an ensemble of three classifiers voting

Two common aprochaes to build ensembles

- **Bagging** (or bootstrap) samples the dataset with replacement
 - The ensemble make prediction by aggregating its predictors
- **Boosting** trains models to correct previous models



(Source)

Ensembles of Decision Trees

Random forests

Trees have poor generalization

- A tree is good doing his job, but does not generalize well
- Different trees could overfit in different ways
- Idea: Using many trees and averaging their result

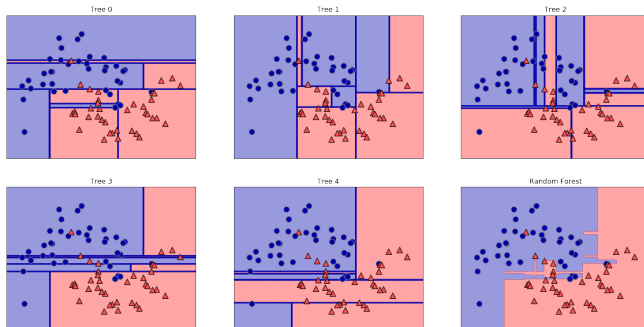
Random forest is an algorithm that trains different trees injecting randomness

- Selecting data - bootstrap
- Selecting features in each test
 - It does not look for the best test
 - It looks for the best test involving a random subset of features
 - The size of the features subset is a critical hyperparameter

Ensembles of Decision Trees

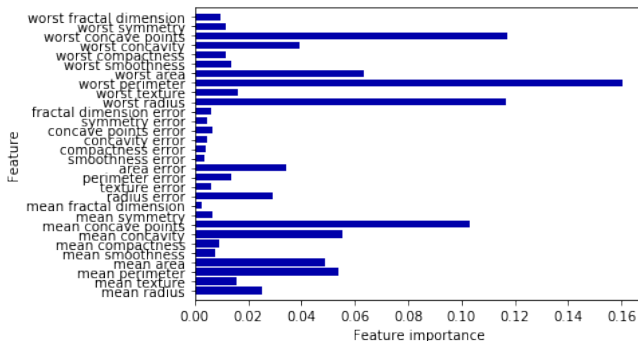
Analyzing random forests (I)

Random forest with five trees



Ensembles of Decision Trees

Analyzing random forests (II)





Feature importance can be aggregated

- More informative than single trees
- The algorithms must consider many possible explanations

Ensembles of Decision Trees

Analyzing random forests (III)



Ensembles of Decision Trees

Gradient boosted regression trees (I)

Random forest is an algorithm that trains different trees

- It encourages diversity by injecting randomness
- Selecting data - bootstrapping
- Selecting features in each test

Ensembles of Decision Trees

Ensembles of Decision Trees : Scikit-learn

`sklearn.ensemble.RandomForestClassifier`

Constructor arguments:

- `n_estimators`: int, default=100
- `max_features`: "sqrt", "log2", None, int or float, default="sqrt"
- `bootstrap`: bool, default=True
- `max_samples`: int or float, default=None Same than `RandomForestClassifier`

Attributes:

- `feature_importances_`: ndarray (n_features,)
- `estimators_`: List of `DecisionTreeClassifier`

Methods: `fit()`, `fit_predict()`

(Scikit-Learn reference)

Ensembles of Decision Trees

Ensembles of Decision Trees : Scikit-learn

```
sklearn.ensemble.RandomForestRegressor
```

Constructor arguments:

- `n_estimators`: int, default=100
- `max_features`: "sqrt", "log2", None, int or float, default=n_features
- `bootstrap`: bool, default=True
- `max_samples`: int or float, default=None Same than RandomForestRegressor

Attributes:

- `feature_importances_`: ndarray (n_features,)
- `estimators_`: List of DecisionTreeRegressor

Methods: `fit()`, `fit_predict()`

(Scikit-Learn reference)

Ensembles of Decision Trees

Summary

Hyperparameters

Advantages

Disadvantages

Support Vector Machines

TODO

Support Vector Machines

Kernelized Support Vector Machines

TODO

Scikit-Learn

Support Vector Machines

Scikit-learn

```
sklearn.cluster.AgglomerativeClustering
```

Constructor arguments:

- `linkage`: 'ward', 'complete', 'average', 'single'

Attributes:

- `n_clusters`: int
- `labels_`: ndarray (n_samples)

Methods: `fit()`, `fit_predict()`

(Scikit-Learn reference)

Support Vector Machines

Summary

Hyperparameters

Advantages

Disadvantages

Algorithms

ARIMA (I)

AR: Autoregressive model

- Current observation depends on the last p observations
- Long term memory

MA: Moving Average model

- Current observation linearly depends on the last q innovations
- Short term memory

ARMA model = AR + MA

- ARMA(p, q): Two hyperparameters, p and q

AR(p)

$$X_t = c + \sum_{i=1}^p \phi_i X_{t-i} + \epsilon_t$$

MA(q)

$$X_t = \mu + \epsilon_t + \theta_1 \epsilon_{t-1} + \dots + \theta_q \epsilon_{t-q}$$


Algorithms

ARIMA (II)

ARIMA = AR + i + MA (AR integrated MA)

- ARIMA(p, d, q)
- Three integer parameters: p, q and d (in practice, low order models)

Forecasts from STL + ARIMA(1,1,1) with drift



(Source)

autoarima: search over p, q and d